

GCD: VLSI's Hello World

EE241 Tutorial

Written by Yunsup Lee (2010)

Updated by Brian Zimmer (2011, 2013)

Overview

For this tutorial, you will become familiar with the VLSI tools you will use throughout this semester and learn how a design “flows” through the toolflow. Specifically, given an RTL model of a simple greatest common divisor (GCD) circuit, you will synthesize and place and route the design, simulate at every stage, and analyze power.

VLSI Toolflow Introduction

Figure 1 shows the toolflow you will be using for the first lab. You will use Synopsys VCS (`vcs`) to *simulate* and *debug* your RTL design. After you get your design right, you will use Synopsys Design Compiler (`dc_shell-xg-t`) to *synthesize* the design. Synthesis is the process of transforming an RTL model into a gate-level netlist. VCS is used again to simulate the synthesized gate-level netlist. After obtaining a working gate-level netlist, you will use Synopsys IC Compiler (`icc_shell`) to *place and route* the design. Placement is the process by which each standard cell is positioned on the chip, while routing involves wiring the cells together using various metal layers. The tools will provide feedback on the performance and area of your design after both synthesis and place and route. The results from place and route are more realistic but require much more time to generate. After place and route, you will generate and simulate the final gate-level netlist using VCS. Finally you will use this gate-level simulation as a final test for correctness and to generate transition counts for every net in the design. Synopsys PrimeTime PX (`pt_shell`) takes these transition counts as input and correlate them with the capacitance values in the final layout to produce estimated power measurements. The diagram below shows how every tools works together.

Prerequisites

As you can easy tell from the diagram, many different tools are needed to take even a simple design from RTL all the way to transistor-level implementation. Each tool is immensely complicated, and many engineers in industry specialize in only one. In order to produce a VLSI design in a single semester we will need to understand a little about every one.

Each tool has a GUI interface. However, most inputs that the tools need are the same for every design iteration and become repetitive to type, so .tcl scripts provide all of the inputs needed. When you use the GUI, in the terminal window you will see the textual equivalent of each click, and these commands can be added to scripts. To keep files organized, each piece of the toolflow has its own build directory and its own Makefile. The Makefile initializes the program and points at the setup scripts. A top-level Makefile runs each program in succession so that ideally, a single command will push an RTL design all of the way through the flow without any repetitive intervention.

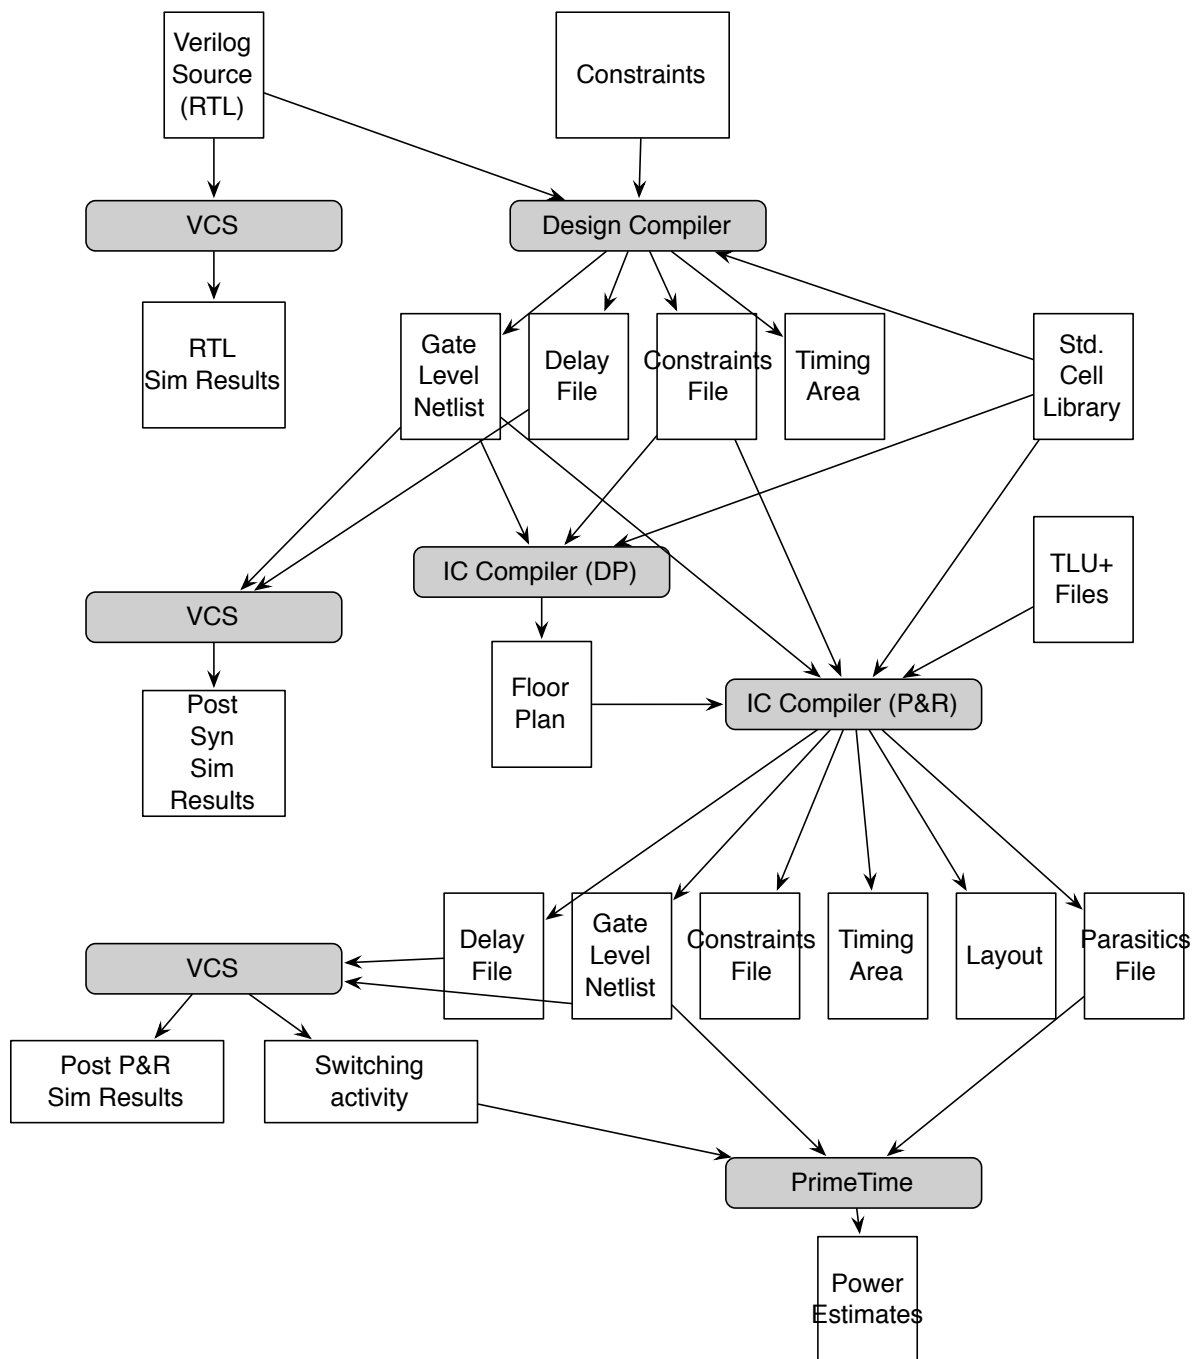


Figure 1: EE241 Toolflow for Lab 1

Getting Started

All of the EE241 laboratory assignments should be completed on an EECS Instructional machine. Please see the course website and follow all of the instructions for setting up your computing resources. Remember, you will need to source a setup script in order for these instructions to work. This bash script contains the location of each tool's binary, and also sets up important environment variables. Make sure you have followed class setup instructions before starting (these are posted on the website).

As these tools generate enormous amounts of data and home directories have too low of a disk quota, we will need to use the local disk of one of the available. Assuming your username is `userA` (change this to your own username), you can create your personal git directory using the following command.

```
% cd /scratch/  
% mkdir userA
```

To begin the lab you will need to make use of a provided lab harness. This lab harness provides makefiles, scripts, and the Verilog test harness required to complete the tutorial. The following commands grab these files from the class repository. To simplify the rest of the lab we will also define a `'$LABROOT'` environment variable which contains the absolute path to the project's top-level root directory. If you're not using bash shell, you can switch to it by typing `bash`.

```
% bash  
% cd /scratch/userA  
% git clone ~ee241/tutorials/gcd  
% cd gcd  
% LABROOT=$PWD
```

Every time you want to start the tools, you must source your environment variable.

```
% source ~ee241/tutorials/ee241.bashrc
```

Note: `scratch/` is a local drive, so if you every need to do work on another machine, you will need to `rsync` files between machines.

The resulting `$LABROOT` directory contains the following primary subdirectories: `src` contains your source Verilog; `build` contains automated makefiles and scripts for building your design. The `src` directory contains the Verilog test harness and other Verilog modules you will need in this lab assignment. Figure 2 shows each directory that you have been given and includes comments about what they do.

RTL:

- `src/gcdGCDUnit_rtl.v` - RTL implementation of `gcdGCDUnit`
- `src/gcdGCDUnitCtrl.v` - Control part of the RTL implementation
- `src/gcdGCDUnitDpath.v` - Datapath part of the RTL implementation
- `src/gcdTestHarness_rtl.v` - Test harness for the RTL model

```

gcd/
  build/    VLSI toolflow for src/
  Makefile/ Controls all pieces of toolflow, e.g. "make dc-syn" will synthesize
  vcs-sim-rtl/ Simulate RTL in ../src/
  dc-syn/    Synthesize RTL in ../src/
  vcs-sim-gl-syn/ Simulate synthesized netlist in dc-syn/current-dc
  icc-par/   Place and route synthesized netlist from dc-syn/current-dc
  vcs-sim-gl-par/ Simulate placed and routed netlist in icc-par/current-icc
  pt-pwr/   Power analysis of design in icc-par/current-icc
  src/      Verilog code

```

Figure 2: Directory organization for lab1-verilog/

The block diagram is shown in Figure 3. Your module is named `gcdGCDUnit` and has the interface shown in Figure 4. We have provided you with a test harness that will drive the inputs and check the outputs of your design.

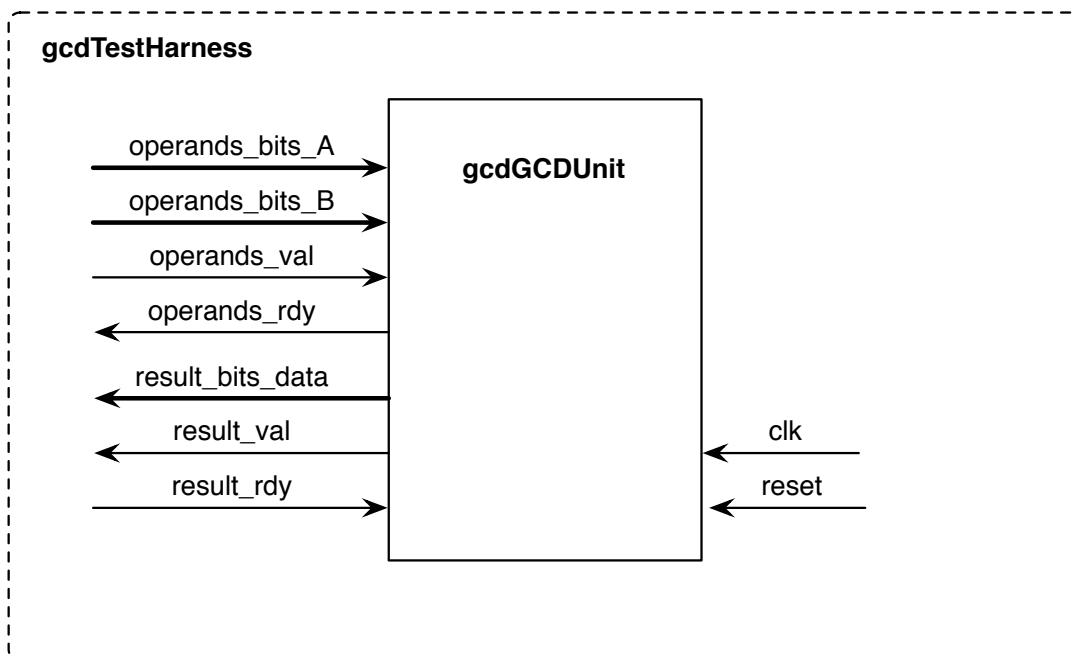


Figure 3: Block diagram for GCD Test Harness

```

module gcdGCDUnit#( parameter W = 16 )
(
    input clk, reset,

    input  [W-1:0] operands_bits_A,    // Operand A
    input  [W-1:0] operands_bits_B,    // Operand B
    input          operands_val,        // Are operands valid?
    output         operands_rdy,        // ready to take operands

    output [W-1:0] result_bits_data,    // GCD
    output         result_val,          // Is the result valid?
    input          result_rdy           // ready to take the result
);

```

Figure 4: Interface for the GCD module

The **build** directory contains the following subdirectories which you will use when building your chip. The order in which they are listed is also the order in which they should be used in the flow.

- **vcs-sim-behav** - Behavioral simulation using Synopsys VCS
- **vcs-sim-rtl** - RTL simulation using Synopsys VCS
- **dc-syn** - Synthesis using Synopsys Design Compiler
- **vcs-sim-gl-syn** - Post synthesis gate-level simulation using Synopsys VCS
- **icc-par** - Automatic placement and routing using Synopsys IC Compiler
- **vcs-sim-gl-par** - Post place and route gate-level simulation using Synopsys VCS
- **pt-pwr** - Power analysis using Synopsys PrimeTime PX

Each subdirectory includes its own makefile and additional script files. So for example, to synthesize with Design Compiler (DC):

```

% cd $LABROOT/build
% cd dc-syn
% make

```

Note: you must follow the ordering given in the list above.

Once you have all the tools working you can use the toplevel makefile in the **build** directory to run multiple tools at once. For example, once all the scripts are properly setup you should be able to use the following command to synthesize, floorplan, and place and route your design. You give the command for the furthest step in the flow you would like to go to, and the Makefile's dependencies ensures that every step before this step is completed first.

```

% cd $LABROOT/build
% make icc-par

```

Makefiles are designed to only rerun when one of their dependencies change. However there are changes you can make that will not trigger a re-run of a step. To force a re-run of any step, you can delete all generated files with 'make clean':

```
% cd $LABROOT/build/dc-syn
% make clean
```

This will delete all previously saved runs in the build- directories. If you would like to keep these, instead of running 'make clean', just delete the current-dc or current-icc directory. Note that for simulation directories vcs-, always use make clean.

Pushing the design through all the VLSI Tools

You will now go through the entire tool flow and inspect the results after each step.

Synopsys VCS: Simulating your Verilog

VCS compiles source Verilog into a cycle-accurate executable for simulation. VCS can compile Verilog expressed behaviorally, at the RTL level, or as structural verilog (a netlist). Behavioral Verilog cannot be synthesized and should only be used in testbenches. RTL-level Verilog expresses behavior as combinational and sequential logic at a higher level. Structural-level Verilog expresses behavior as specific gates wired together. You will start with simulating the GCD module RTL (before it is synthesized).

```
% cd $LABROOT/build/vcs-sim-rtl
% make
% make run
./simv +verbose=1
...
+ Running Test Case: gcdGCDUnit_rtl
  [ passed ] Test ( vcTestSink ) succeeded, [ 0003 == 0003 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0007 == 0007 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0005 == 0005 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0001 == 0001 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0028 == 0028 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 000a == 000a ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0005 == 0005 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0000 == 0000 ]
  [ passed ] Test ( Is sink finished? ) succeeded
...
```

Where should you start if all of your tests didn't pass? The answer is debug your RTL using Discovery Visualization Environment (DVE) GUI looking at the trace outputs. The simulator already logged the activity for every net to the `vcdplus.vpd` file. DVE can read the `vcdplus.vpd` file and visualize the wave form.

```
% dve -vpd vcdplus.vpd &
```

To add signals to the waveform window (see Figure 5) you can select them in the hierarchy window and then right click to choose *Add To Waves > New Wave View*.

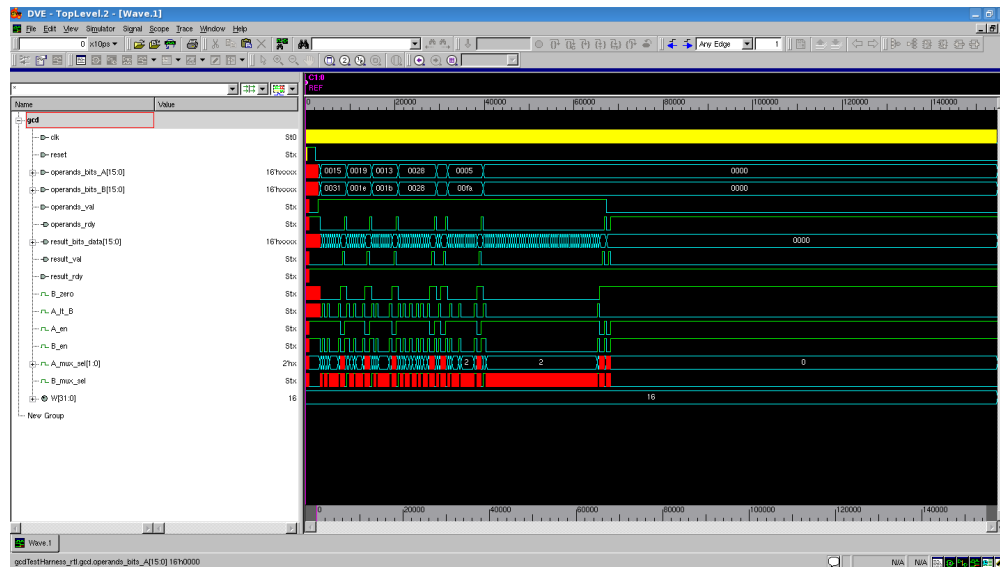


Figure 5: DVE Waveform Window

Synopsys Design Compiler: RTL to Gate-Level Netlist

Design Compiler performs hardware synthesis. A synthesis tool takes an RTL hardware description and a standard cell library as input and produces a gate-level netlist as an output. The resulting gate-level netlist is a completely structural description with only standard cells at the leaves of the design.

```
% cd $LABROOT/build/dc-syn
% make
```

Go ahead and take a look what the automated build system produced.

```
% cd $LABROOT/build/dc-syn
% ls -l
-rw-r--r-- 1 yunsup grad 4555 Aug 29 22:15 Makefile
drwxr-xr-x 7 yunsup grad 4096 Aug 29 22:15 build-dc-2010-08-29_22-15
-rw-r--r-- 1 yunsup grad 1108 Aug 28 12:06 constraints.tcl
lrwxrwxrwx 1 yunsup grad 25 Aug 29 22:15 current-dc -> build-dc-2010-08-29_22-15
drwxr-xr-x 2 yunsup grad 4096 Aug 29 21:39 rm_dc_scripts
drwxr-xr-x 2 yunsup grad 4096 Aug 28 12:00 rm_notes
drwxr-xr-x 2 yunsup grad 4096 Aug 29 21:50 rm_setup
% cd current-dc
% ls -l
drwxr-xr-x 2 yunsup grad 4096 Aug 29 22:15 WORK
-rw-r--r-- 1 yunsup grad 47 Aug 29 22:15 access.tab
-rw-r--r-- 1 yunsup grad 235827 Aug 29 22:15 command.log
-rw-r--r-- 1 yunsup grad 5141 Aug 29 22:15 common_setup.tcl
-rw-r--r-- 1 yunsup grad 1108 Aug 29 22:15 constraints.tcl
-rw-r--r-- 1 yunsup grad 18996 Aug 29 22:15 dc.tcl
```

```

-rw-r--r-- 1 yunsup grad 4621 Aug 29 22:15 dc_setup.tcl
-rw-r--r-- 1 yunsup grad 4625 Aug 29 22:15 dc_setup_filenames.tcl
-rw-r--r-- 1 yunsup grad 2730 Aug 29 22:15 find_regs.tcl
-rw-r--r-- 1 yunsup grad 4439 Aug 29 22:15 force_regs.ucli
drwxr-xr-x 3 yunsup grad 4096 Aug 29 22:15 gcdGCDUnit_rtl_LIB
drwxr-xr-x 2 yunsup grad 4096 Aug 29 22:15 log
-rw-r--r-- 1 yunsup grad 1087 Aug 29 22:15 make_generated_vars.tcl
drwxr-xr-x 2 yunsup grad 4096 Aug 29 22:15 reports
drwxr-xr-x 2 yunsup grad 4096 Aug 29 22:15 results
-rw-r--r-- 1 yunsup grad 29 Aug 29 22:15 timestamp

```

Notice that the makefile does not overwrite build directories. It always create new build directories. This makes it easy to change your synthesis scripts or source Verilog, resynthesize your design, and compare your results to previous designs. You can use symlinks to keep track of various build directories. Inside the `current-dc` directory, you can see all the tcl scripts as well as the directories named `results` and `reports`: `results` contains your synthesized gate-level netlist; and `reports` contains various post synthesis reports.

Take a look at various reports on synthesis results.

```

% cd $LABROOT/build/dc-syn/current-dc/reports
% cat gcdGCDUnit_rtl.mapped.timing.rpt
Startpoint: GCDdpath0/B_reg_reg_11_
(rising edge-triggered flip-flop clocked by ideal_clock1)
Endpoint: GCDdpath0/A_reg_reg_5_
(rising edge-triggered flip-flop clocked by ideal_clock1)
Path Group: ideal_clock1
Path Type: max

```

Point	Fanout	Cap	Trans	Incr	Path
<hr/>					
clock ideal_clock1 (rise edge)				0.0000	0.0000
clock network delay (ideal)				0.0000	0.0000
GCDdpath0/B_reg_reg_11_/CLK (DFFARX1_RVT)			0.0000	0.0000	0.0000 r
GCDdpath0/B_reg_reg_11_/Q (DFFARX1_RVT)			0.0269	0.0999	0.0999 f
GCDdpath0/B_reg[11] (net)	2	2.4320		0.0000	0.0999 f
GCDdpath0/U77/A (INVX2_RVT)			0.0269	0.0000 *	0.0999 f
GCDdpath0/U77/Y (INVX2_RVT)			0.0182	0.0167	0.1166 r
GCDdpath0/n38 (net)	3	2.4349		0.0000	0.1166 r
GCDdpath0/U291/A1 (NAND2X0_RVT)			0.0182	0.0000 *	0.1167 r
GCDdpath0/U291/Y (NAND2X0_RVT)			0.0296	0.0238	0.1405 f
GCDdpath0/n249 (net)	2	1.1267		0.0000	0.1405 f
GCDdpath0/U284/A1 (NAND2X0_RVT)			0.0296	0.0000 *	0.1405 f
GCDdpath0/U284/Y (NAND2X0_RVT)			0.0370	0.0375	0.1780 r
GCDdpath0/n75 (net)	1	1.7944		0.0000	0.1780 r
GCDdpath0/U74/A1 (A021X1_RVT)			0.0370	0.0000 *	0.1780 r
GCDdpath0/U74/Y (A021X1_RVT)			0.0193	0.0454	0.2233 r
GCDdpath0/n36 (net)	1	0.5584		0.0000	0.2233 r
GCDdpath0/U149/A3 (A022X1_RVT)			0.0193	0.0000 *	0.2233 r
GCDdpath0/U149/Y (A022X1_RVT)			0.0235	0.0390	0.2623 r
GCDdpath0/n295 (net)	1	1.4197		0.0000	0.2623 r
GCDdpath0/U229/A (INVX2_RVT)			0.0235	0.0000 *	0.2623 r
GCDdpath0/U229/Y (INVX2_RVT)			0.0148	0.0104	0.2727 f

GCDdpath0/n1 (net)	1	1.3086		0.0000	0.2727 f
GCDdpath0/U107/A3 (OA21X1_RVT)			0.0148	0.0000 *	0.2728 f
GCDdpath0/U107/Y (OA21X1_RVT)			0.0230	0.0381	0.3109 f
GCDdpath0/A_lt_B (net)	3	2.0316		0.0000	0.3109 f
GCDdpath0/A_lt_B (gcdGCDUnitDpath_W16)				0.0000	0.3109 f
A_lt_B (net)		2.0316		0.0000	0.3109 f
GCDctrl0/A_lt_B (gcdGCDUnitCtrl)				0.0000	0.3109 f
GCDctrl0/A_lt_B (net)		2.0316		0.0000	0.3109 f
GCDctrl0/U7/A (INVX1_RVT)			0.0230	0.0000 *	0.3109 f
GCDctrl0/U7/Y (INVX1_RVT)			0.0171	0.0166	0.3275 r
GCDctrl0/n18 (net)	2	1.2895		0.0000	0.3275 r
GCDctrl0/U15/A1 (AND2X1_RVT)			0.0171	0.0000 *	0.3275 r
GCDctrl0/U15/Y (AND2X1_RVT)			0.0214	0.0344	0.3619 r
GCDctrl0/A_mux_sel[1] (net)	2	2.2309		0.0000	0.3619 r
GCDctrl0/A_mux_sel[1] (gcdGCDUnitCtrl)				0.0000	0.3619 r
A_mux_sel[1] (net)		2.2309		0.0000	0.3619 r
GCDdpath0/A_mux_sel[1] (gcdGCDUnitDpath_W16)				0.0000	0.3619 r
GCDdpath0/A_mux_sel[1] (net)		2.2309		0.0000	0.3619 r
GCDdpath0/U294/A (NBUFFX2_RVT)			0.0214	0.0000 *	0.3619 r
GCDdpath0/U294/Y (NBUFFX2_RVT)			0.0291	0.0397	0.4015 r
GCDdpath0/n72 (net)	8	7.5054		0.0000	0.4015 r
GCDdpath0/U41/A1 (AND2X1_RVT)			0.0291	0.0001 *	0.4016 r
GCDdpath0/U41/Y (AND2X1_RVT)			0.0158	0.0301	0.4317 r
GCDdpath0/n119 (net)	1	0.7624		0.0000	0.4317 r
GCDdpath0/U110/A1 (A022X1_RVT)			0.0158	0.0000 *	0.4317 r
GCDdpath0/U110/Y (A022X1_RVT)			0.0225	0.0470	0.4788 r
GCDdpath0/A_next[5] (net)	1	1.1741		0.0000	0.4788 r
GCDdpath0/A_reg_reg_5_/D (DFFARX1_RVT)			0.0225	0.0000 *	0.4788 r
data arrival time					0.4788
<hr/>					
clock ideal_clock1 (rise edge)				0.5000	0.5000
clock network delay (ideal)				0.0000	0.5000
clock uncertainty				-0.0250	0.4750
GCDdpath0/A_reg_reg_5_/CLK (DFFARX1_RVT)				0.0000	0.4750 r
library setup time				-0.0315	0.4435
data required time					0.4435
<hr/>					
data required time					0.4435
data arrival time					-0.4788
<hr/>					
slack (VIOLATED)					-0.0353

...

This report lists the *critical path* of the design. The critical path is the slowest logic between any two registers and is therefore the limiting factor preventing you from decreasing the clock period constraint. In the example above (which is not necessarily what you'll see in your tutorial), you can see that the critical path starts at bit 11 of the operand B register in the datapath; goes through the comparator; to the control logic; and finally ends bit 5 of operand A register in the datapath. The critical path takes a total of 0.4788ns which is less than the 0.5ns clock period constraint. However, your timing constraint is still VIOLATED due to accounting for clock skew and setup time.

If any of your paths say (VIOLATED) instead of (MET), you need to modify the `build/Makefrag` file to increase the clock period. If you do not do this, your design will not pass simulation. Note also: if you make your clock period too tight for a given design, the tools (synthesis and place and route) will have trouble doing their jobs and the run time will increase!

```
% vim $LABROOT/build/Makefrag
...
clock_period = 0.8
...
% cd $LABROOT/build/dc-syn
% make clean
% make
% cd $LABROOT/build/dc-syn/current-dc/reports
% cat gcdGCDUnit_rtl.mapped.timing.rpt
```

Double check that timing is met now. Additionally, you can get a sense of whether something went horribly wrong or not by looking at the QOR report, which summarizes things like area and timing.

```
% cat gcdGCDUnit_rtl.mapped.qor.rpt
```

As you're pushing your design through the tools for the first time, you'll likely be making adjustments to the synthesis and place and route scripts provided to you. Sometimes, the changes might cause errors, so it's always important to keep a careful eye on the log files for errors and warnings. Word of caution: the tools generate a lot of benign warning messages too, so you'll need to get a feel of which messages matter.

```
% cd $LABROOT/build/dc-syn/current-dc/log
% grep Error *
% grep Warn* *

% cd $LABROOT/build/dc-syn/current-dc/reports
% cat gcdGCDUnit_rtl.mapped.area.rpt
...
```

	Global cell area		Local cell area		
	Absolute	Percent	Combi-	Noncombi-	Black
	Total	Total	national	national	boxes
Hierarchical cell					
-----	-----	-----	-----	-----	-----
gcdGCDUnit_rtl	966.5096	100.0	11.6906	0.0000	0.0000
GCDctrl0	56.4200	5.8	43.2045	13.2155	0.0000
GCDdpath0	898.3991	93.0	658.4871	228.2213	0.0000
GCDdpath0/clk_gate_A_reg_reg	5.8453	0.6	0.0000	5.8453	0.0000
GCDdpath0/clk_gate_B_reg_reg	5.8453	0.6	0.0000	5.8453	0.0000
-----	-----	-----	-----	-----	-----
Total			713.3822	253.1274	0.0000
...					

This report tells you the post synthesis area results. The units are μm^2 . You can see that the datapath consumes 93.3% of the total chip area.

```
% cd $LABROOT/build/dc-syn/current-dc/reports
% cat gcdGCDUnit_rtl.mapped.power.rpt
...
```

Hierarchy	Switch Power	Int Power	Leak Power	Total Power	%
gcdGCDUnit_rtl	167.522	467.149	1.25e+08	760.139	100.0
GCDctrl0 (gcdGCDUnitCtrl)	5.927	35.039	6.75e+06	47.713	6.3
GCDdpath0 (gcdGCDUnitDpath_W16)	154.322	427.257	1.13e+08	695.001	91.4
...					

This report tells you about post synthesis power results using some default activity factor. The dynamic power units are *uW* while the leakage power units are *pW*.

```
% cd $LABROOT/build/dc-syn/current-dc/reports
% cat gcdGCDUnit_rtl.mapped.reference.rpt
...
```

Design: gcdGCDUnitCtrl

Reference	Library	Unit Area	Count	Total Area	Attributes
AND2X1_RVT	saed32rvt_tt1p05v25c	2.033152	3	6.099456	
AND2X2_RVT	saed32rvt_tt1p05v25c	2.287296	3	6.861888	
DFFX1_RVT	saed32rvt_tt1p05v25c	6.607744	2	13.215488	n
INVX1_RVT	saed32rvt_tt1p05v25c	1.270720	3	3.812160	
NAND2X0_RVT	saed32rvt_tt1p05v25c	1.524864	1	1.524864	
NAND3X0_RVT	saed32rvt_tt1p05v25c	1.779008	1	1.779008	
NAND4X0_RVT	saed32rvt_tt1p05v25c	2.033152	1	2.033152	
NBUFFX2_RVT	saed32rvt_tt1p05v25c	2.033152	1	2.033152	
NOR2X0_RVT	saed32rvt_tt1p05v25c	2.541440	3	7.624320	
OA22X2_RVT	saed32rvt_tt1p05v25c	2.795584	1	2.795584	
OR2X1_RVT	saed32rvt_tt1p05v25c	2.033152	2	4.066304	
OR2X2_RVT	saed32rvt_tt1p05v25c	2.287296	2	4.574592	

Total 12 references 56.419969

...

This report lists the standard cells used in each module.

```
% cd $LABROOT/build/dc-syn/current-dc/reports
% cat gcdGCDUnit_rtl.mapped.resources.rpt
...
```

Cell	Module	Current Implementation	Set Implementation
sub_x_2	DW01_sub	pparch (area,speed)	

```
| lt_x_3          | DW_cmp          | apparch (area)    |
=====
```

...

Synopsys provides a library of commonly used arithmetic components as highly optimized building blocks. This library is called Design Ware and Design Compiler will automatically use Design Ware components when it can. This report can help you determine when Design Compiler is using Design Ware components. The DW01_sub in the module name indicates that this is a Design Ware subtractor. This report also gives you what type of architecture it used.

Last, see what you actually created by viewing the post-synthesis netlist.

```
% cd $LABROOT/build/dc-syn/current-dc/results
% cat gcdGCDUnit_rtl.mapped.v
```

Synopsys provides a GUI front-end for Design Compiler called Design Vision which you will use to analyze the synthesis result. You should avoid using the GUI to actually perform synthesis since you want to use scripts for this. Now launch design vision.

```
% cd $LABROOT/build/dc-syn/current-dc
% ./start_gui
```

You can browse your design with the hierarchical view (see Figure 6). If you right click on a module and choose *Schematic View* option, the tool will display a schematic of the synthesized logic corresponding to that module.

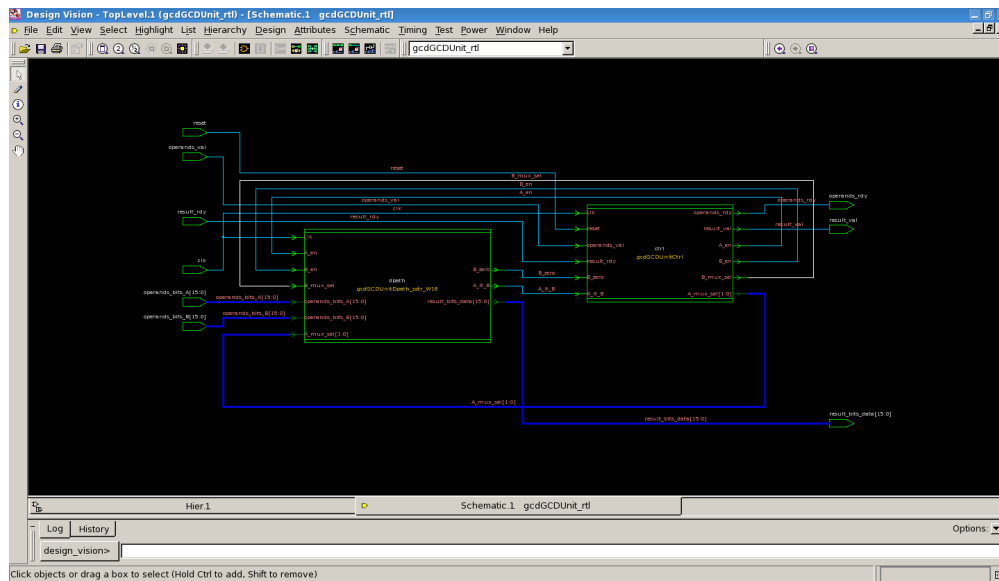


Figure 6: Design Vision Hierarchical View

Synopsys VCS: Simulating Post Synthesis Gate-Level Netlist

After obtaining the synthesized gate-level netlist, you will double-check the netlist by running a simulation using VCS.

```
% cd $LABROOT/build/vcs-sim-gl-syn
% make
% make run
...
+ Running Test Case: gcdGCDUnit_rtl
  [ passed ] Test ( vcTestSink ) succeeded, [ 0003 == 0003 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0007 == 0007 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0005 == 0005 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0001 == 0001 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0028 == 0028 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 000a == 000a ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0005 == 0005 ]
  [ passed ] Test ( vcTestSink ) succeeded, [ 0000 == 0000 ]
  [ passed ] Test ( Is sink finished? ) succeeded
...
```

Synopsys IC Compiler: Gate-Level Netlist to Layout

IC Compiler performs place and route. This tool takes a synthesized gate-level netlist and a standard cell library as input and produces a layout as an output.

One critical part of the place and route process is coming up with the initial floorplan. Take a look at the floorplan script used just to get a feel for things. In your floorplan script, you tell the tools where to place macros like SRAMs and how to place power supply straps (spacing, direction), among other things. To get a design with good QOR, you'll probably be spending a considerable amount of time floorplanning. Doing a good job telling the tool how to do things also helps with minimizing DRC violations.

```
% cat $LABROOT/build/icc-par/floorplan/floorplan.tcl
```

You can automate the process of running ICC. Notice that the makefile creates new build directories like the one in Design Compiler.

```
% cd $LABROOT/build/icc-par
% make
% ls -l
-rw-r--r-- 1 yunsup grad 15232 Aug 28 18:40 Makefile
drwxr-xr-x 6 yunsup grad 4096 Aug 29 23:42 build-icc-2010-08-29_23-41
drwxr-xr-x 8 yunsup grad 4096 Aug 29 23:41 build-iccdp-2010-08-29_23-41
lrwxrwxrwx 1 yunsup grad 26 Aug 29 23:41 current-icc -> build-icc-2010-08-29_23-41
lrwxrwxrwx 1 yunsup grad 28 Aug 29 23:41 current-iccdp -> build-iccdp-2010-08-29_23-41
drwxr-xr-x 2 yunsup grad 4096 Aug 29 23:19 rm_icc_dp_scripts
drwxr-xr-x 2 yunsup grad 4096 Aug 29 23:29 rm_icc_scripts
drwxr-xr-x 2 yunsup grad 4096 Aug 29 23:38 rm_icc_zrt_scripts
drwxr-xr-x 2 yunsup grad 4096 Aug 28 16:20 rm_notes
drwxr-xr-x 2 yunsup grad 4096 Aug 29 23:34 rm_setup
```

After a few minutes, routing should have finished. Browse the results directory files to see similar outputs as Design Compiler. Again, it's good to sanity-check the QOR output and logs (both for floorplan in current-iccdp and for the rest of the P&R steps in current-icc) for warnings and errors.

```
% cd $LABROOT/build/icc-par/current-iccdp/log
% grep Error *
% grep Warn* *
% cd $LABROOT/build/icc-par/current-icc/log
% grep Error *
% grep Warn* *
% cat $LABROOT/build/icc-par/current-icc/reports/chip_finish_icc.qor.rpt
```

Now open the GUI:

```
% cd $LABROOT/build/icc-par/current-icc
% ./start_gui
```

Take a look at the generated clock tree. Choose *Clock > Color By Clock Trees*. Hit *Reload*, and then hit *OK* on the popup window. Now you will be able to see the synthesized clock tree (Figure 7).




Figure 7: Synthesized clock tree shown in IC Compiler

Figure 8 shows the routed signals. Synopsys 32nm process provides nine metal layers (metal 1 is mostly used by the standard cell layout itself) to route your signals.

IC Compiler actually performs place and route in different steps. The Makefile runs then all at once by default (That's why you see it starting/stopping if you're following the terminal output!), but if you're interacting with the tool directly, you can run only certain steps at a time. Go to File — Open Design, then try opening all of the different steps (for example: `place_opt.icc`), which shows ICC's first placement.

The post place-and-route netlist can be found as `current-icc/results/gcdGCDUnit_rtl.output.v`.

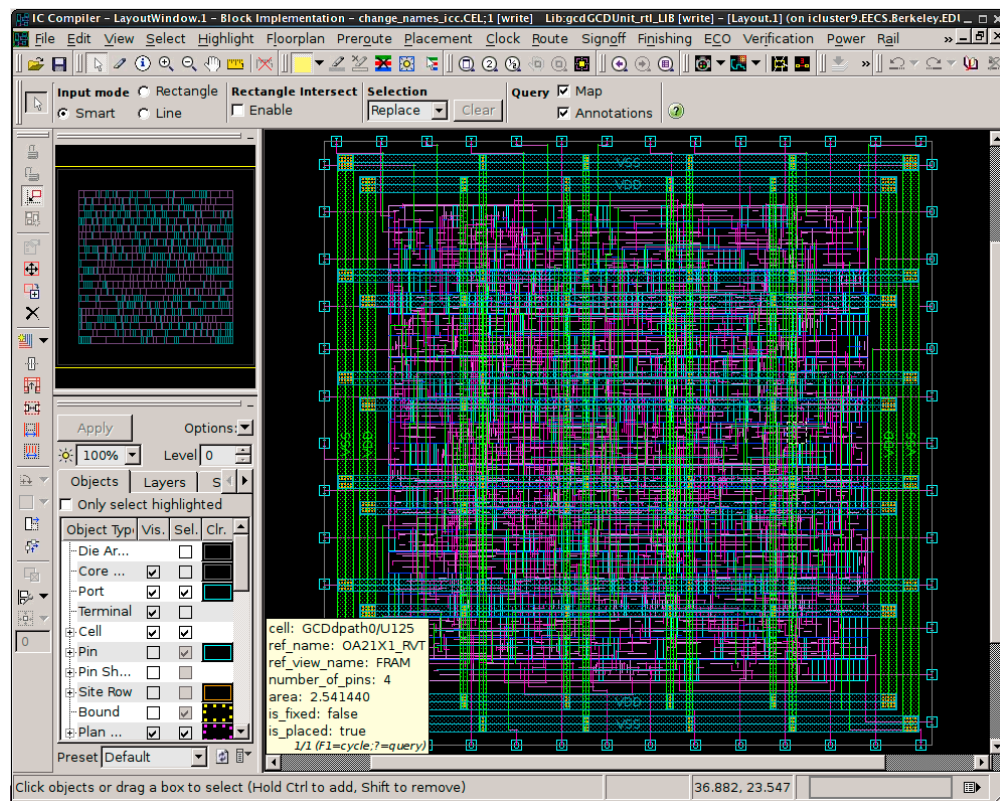


Figure 8: Routed signals shown in IC Compiler

Synopsys VCS: Simulating Post Place and Route Gate-Level Netlist

After you obtain the post place and route gate-level netlist, you will double-check the netlist by running a simulation using VCS. You can use the makefile to build the post synthesis gate-level netlist simulator, run, and convert the switching activity file into a vcd and a saif format (used by PrimeTime later).

```
% cd $LABROOT/build/vcs-sim-gl-par
% make
% make run
% make convert
```

If your circuit no longer works (and it did after synthesis), this means that the added parasitics from routing and clock tree issues increased your critical path. You can increase the clock period for the simulation testbench by editing Makefrag and then rerun the above commands (although, for this design, you shouldn't have to if you changed the clock period to a fairly conservative 0.8ns previously).

```
% vim $LABROOT/build/Makefrag
...
vcs_clock_period = 0$(shell echo "scale=4; ${clock_period}*0.5*1.2" | bc)
...
```

Synopsys PrimeTime PX: Estimating Power

PrimeTime PX is an add-on feature to PrimeTime that analyzes power dissipation of a cell-based design. PrimeTime PX supports two types of power analysis modes. They are averaged mode and time-based mode. Averaged mode calculates averaged power based on toggle rates. Time-based mode let's you know the peak power as well as the averaged power using gate-level simulation activity.

```
% cd $LABROOT/build/pt-pwr
% make
```

Now look at some of the results

```
% cd $LABROOT/build/pt-pwr/current-pt/reports
% cat vcdplus.power.avg.max.report
...
```

Hierarchy	Switch Power	Int Power	Leak Power	Total Power	%
gcdcUnit_rtl	3.84e-05	1.34e-04	1.57e-04	3.29e-04	100.0
GCDctrl0 (gcdcUnitCtrl)	2.60e-05	9.36e-05	1.24e-05	1.32e-04	40.2
GCDdpath0 (gcdcUnitDpath_W16)	1.22e-05	3.99e-05	1.41e-04	1.93e-04	58.8

```
% cat vcdplus.power.time.max.report
```

Hierarchy	Switch Power	Int Power	Leak Power	Peak Power	Peak Time
-----------	-----------------	--------------	---------------	---------------	--------------

gcdGCDUnit_rtl	3.84e-05	1.33e-04	1.57e-04	9.62e-03	49.331
GCDctrl0 (gcdGCDUnitCtrl)	2.60e-05	9.35e-05	1.24e-05	1.82e-03	43.914
GCDdpath0 (gcdGCDUnitDpath_W16)	1.22e-05	3.97e-05	1.41e-04	8.81e-03	49.331

Summary

You have now taken an RTL-level description in Verilog, then synthesized, place-and-routed, and analyzed the power of the design. Each step is automated through Makefiles, and each step has its own directory.

If you are interesting in learning more about how these scripts run, open the .log files in the directory to see the commands that are run in each tool. Then open the tool directly (eg. type `dc_shell`) and paste in each line.

HAMMER

So far we've pushed a GCD block through an entire digital design flow using tools provided by Synopsys. However, all the digital flow you've just run has been written beforehand and was tailored to operating on the Synopsys toolset. Running this design on another set of tools, such as the design suite offered by Cadence, would require the scripts rewritten with commands for the new tools. Similarly, redoing the design in a new technology would require redefining all the library lists.

HAMMER is a tool that aims to codify the design flow process and preserve just the methodology of designing a chip, while separating other issues such as tool-specific scripts and technology files so that ideally you could write a HAMMER flow for one technology and toolset, and simply swap out the tools or technology when necessary.

Getting Started

First, clone the git repo containing the hammer setup and go to the `lab1` branch.

```
% git clone ~ee241/hammer-stuff/ee241bS2019
% cd ee241bS2019
% git checkout lab1
% git submodule update --init --recursive
% source sourceme.sh
```

HAMMER exists as an open-source repository on GitHub, and it is included in the repo we just cloned as a submodule. Running the update automatically clones these subrepos and now we should be ready to go to begin our flow!

0.1 Synthesis

This HAMMER flow is set up to use the Cadence digital design tools rather than the Synopsys ones you've used before. The Cadence synthesis tool is called Genus and for what we care about today, it is the same as if we were putting the design through DC. To begin running synthesis, type in this command into the terminal.

```
% hammer-vlsi synthesis -p gcdGCDUnit_rtl.yml -o syn-output.json
```

Let's see what this command is doing. `hammer-vlsi` is the driver behind the HAMMER flow which executes the generation steps. HAMMER is composed of steps, each performing some action relevant to the flow, such as performing clock tree synthesis or routing optimization. The `-p` option tells the driver that it should expect an input file of `gcdGCDUnit_rtl.yml`, which describes the Verilog files it needs, and the `-o` option is an output file that remembers paths to other files that we'll need for the next step.

Place and Route

Cadence's P&R tool is called Innovus, and we will be using this tool next. To begin place and route, we must first prepare some files for Innovus to use. Type the following command into the terminal.

```
% hammer-vlsi syn_to_par -p gcdGCDUnit_rtl.yml -p syn-output.json -o par-run.json
```

This is a quick step and generates a new JSON file that will let HAMMER know some paths and constraints to generate for Innovus. Now we get to the fun part! Run the following command.

```
% hammer-vlsi par -p par-run.json
```

This runs the place and route sequence where HAMMER now generates scripts and constraint files as specified in the YAML file before. Let this run finish and take a look at the final result. The outputs of the syn and par steps are in `obj/syn-rundir` and `obj/par-rundir` by default. Look in there for the output log, (called `innovus.log*`). Open the file in your favorite text editor and do a search for "opt_design Final Summary" to see how much of the chip area we are using.

```
-----
opt_design Final Summary
-----
```

Setup views included:

my_view

Setup mode	all	reg2reg	default
WNS (ns):	0.000	N/A	0.000
TNS (ns):	0.000	N/A	0.000
Violating Paths:	0	N/A	0
All Paths:	0	N/A	0

Real		Total	
DRVs	Nr nets(terms)	Worst Vio	Nr nets(terms)
max_cap	3 (3)	-0.002	3 (3)
max_tran	0 (0)	0.000	0 (0)
max_fanout	0 (0)	0	0 (0)
max_length	0 (0)	0	0 (0)

Density: 49.491%

Routing Overflow: 0.00% H and 0.00% V

49.5% doesn't seem like a very efficient use of chip area, so let's modify the run a little bit. Open up gcdGCDUnit_rtl.yml and take a look at what it's saying.

```
synthesis.inputs.input_files: ["src/lab1_gcd/gcdGCDUnit_rtl.v", "src/lab1_gcd/gcdGCDUnitCtrl
synthesis.inputs.top_module: "gcdGCDUnit_rtl"
#vlsi.inputs.clocks:
#- name: "clock"
  #period: "50 ns"
  #uncertainty: "1 ns"
vlsi.inputs.placement_constraints:
- path: "gcdGCDUnit_rtl"
  type: "toplevel"
  x: 0
  y: 0
  width: 30
  height: 30
```

```

margins:
  left: 0
  right: 0
  top: 0
  bottom: 0
#par.innovus.power_straps_mode: generate

```

Here we see the input Verilog files, the top module name, and the placement constraints we are passing to Innovus. We've given it a $20\mu\text{m} \times 30\mu\text{m}$ area to work with and it's clearly more than enough, so let's try tightening the constraints a bit. We know that at 30×30 it was using only about 50% area, so it should be able to handle half the amount we are allowing it. A quick calculation shows that a square of 21.2×21.2 would give us about half the area of the previous constraint so change x and y and try running the entire flow again, from synthesis to par.

```
**ERROR: (IMPSP-190): Design has util 151.4% > 100%, placer cannot proceed.
```

Looks like that was too strict of a constraint and now Innovus has run out of area to place the design. We can relax the design a bit to something a bit more reasonable, such as 26×26 .

```

-----
opt_design Final Summary
-----

```

Setup views included:

my_view

Setup mode	all	reg2reg	default
WNS (ns):	0.000	N/A	0.000
TNS (ns):	0.000	N/A	0.000
Violating Paths:	0	N/A	0
All Paths:	0	N/A	0

DRVs	Real		Total
	Nr nets(terms)	Worst Vio	Nr nets(terms)
max_cap	12 (12)	-0.052	12 (12)
max_tran	1 (1)	-0.480	1 (1)
max_fanout	0 (0)	0	0 (0)
max_length	0 (0)	0	0 (0)

Density: 97.895%

Routing Overflow: 0.00% H and 0.00% V

Now this is much closer to 100% usage of the chip area. Notice this entire time we haven't had to look at any Tcl scripts or write any Makefiles. HAMMER has abstracted the parts of the flow specific to the new tools we are using to run the same design flow as before, but now using a tool we were completely unfamiliar with.

Acknowledgements

Many people have contributed to versions of this lab over the years. The lab was originally developed for CS250 VLSI Systems Design course at University of California at Berkeley by Yunsup Lee. Contributors include: Krste Asanović, Christopher Batten, John Lazzaro, and John Wawrzynek. Versions of this lab have been used in the following courses:

- CS250 VLSI Systems Design (2009-2010) - University of California at Berkeley
- CSE291 Manycore System Design (2009) - University of California at San Diego